

CSE 6220 PA 1

Karl Hiner, Harshit Srivastava, Dalton Yu

Below is a chart showing the run-time of the program (in seconds) vs. the number of processors for $p = [1, 2, \dots, 24]$. This experiment was performed on the COC-ICE cluster.




Figure 1: $n = 10^6$, $p = [1, 2, \dots, 24]$ on PACE cluster

Note that this chart evaluates every value of p in the tested range. The execution time for $p = 2$ is very nearly half that of $p = 1$, showing a near-ideal speedup. Likewise, $p = 4$ shows roughly a 2X speedup over $p = 2$, as expected for this embarrassingly parallel algorithm.

However, these speedups gradually decrease, reaching a minimum somewhere in the range $p \in [14, 20]$, and execution times even begin to consistently climb back up as $p > 20$. This is because the communication overhead (such as

latency and bandwidth) starts to overwhelm the speedup gained from parallel computation.

Note that with a problem size of $n = 10^6$, using $p = 20$ processors results in each process being responsible for $\frac{10^6}{20} = 5 \cdot 10^4$ summation terms. Calculating each term involves three multiplications and two additions for a total of five operations each. Thus, we see communication overhead dominating the runtime when each processor is responsible for roughly $\leq \frac{1}{4}$ M floating point operations.

Note that we cap our processor count at $p = 24$ since this experiment was run on the COC-ICE cluster, which has 24 cores per node. Thus, after crossing $p = 24$, we encounter dramatically higher execution times, due to inter-*node* communication, which is much more costly than inter-*core* communication.

This is clear from the following chart, which shows $p = [1, 2, \dots, 48]$:




Figure 2: $n = 10^6$, $p = [1, 2, \dots, 48]$ on PACE cluster